

Function Call Example 5 — Recursion:

- Functions can call themselves and/or other functions.
- Set up the stack in the same way as for the original function call
 - Not really any different from what we have been doing!

- Fibonacci sequence:

$$f(N) = \begin{cases} 1, N = 1 \\ 1, N = 2 \\ f(N-1) + f(N-2), N \geq 3 \end{cases}$$

- Need to check the two base cases:

```
# if N == 1, return 1
addi $t0, $zero, 1
bne $a0, $t0, N2
addi $v0, $zero, 1
j    fibend
```

```
N2: # if N == 2, return 1
addi $t0, $zero, 2
bne $a0, $t0, N3
addi $v0, $zero, 1
j    fibend
```

```
N3: # compute fibonacci(N-1) + fibonacci(N-2)
```

Function Call Example 5 — Recursion (continued):

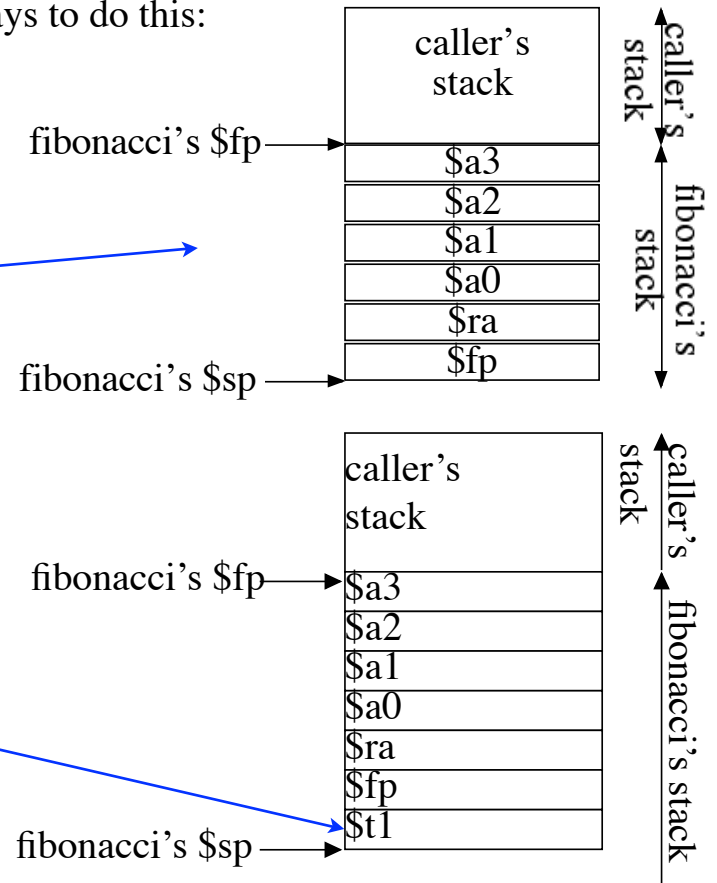
- Need to make two recursive calls:
 - Need to “remember” the value of `$a0` so we can restore it — we have done this part before.
 - Need to “remember” the results of the two recursive calls. Two ways to do this:
 - Use a register for each — `$t1` and `$t2` in my example.
 - Save them as “local” variables.

- Using two registers:

```

N3:      # compute $t1 = fibonacci(N-1)
        addi    $a0, $a0, -1      # compute N - 1
        jal     fibonacci
        add     $t1, $v0, $zero   # save result1 in $t1

        # compute $t2 = fibonacci(N-2)
        lw      $a0, 8($sp)
        # save $t1 on the stack first
        # grow stack temporarily
        addiu   $sp, $sp, -4
        sw      $t1, 0($sp)
        addi    $a0, $a0, -2      # compute N - 2
        jal     fibonacci
        add     $t2, $v0, $zero   # save result2 in $t2
        # get $t1 off the stack and shrink the stack
        lw      $t1, 0($sp)
        addiu   $sp, $sp, 4
    
```



Function Call Example 5 — Recursion (continued):

- The code for fibonacci using registers:

```
fibonacci:
    # Prologue: set up stack and frame pointers for fibonacci
    # Standard 24-byte stack
    addiu    $sp, $sp, -24    # allocate stack space -- default of 24 here
    sw       $fp, 0($sp)     # save caller's frame pointer
    sw       $ra, 4($sp)     # save return address
    sw       $a0, 8($sp)     # save parameter value
    addi     $fp, $sp, 20    # setup fibonacci's frame pointer

    # if N == 1, return 1
    addi     $t0, $zero, 1
    bne      $a0, $t0, N2
    addi     $v0, $zero, 1
    j        fibend

N2:        # if N == 2, return 1
    addi     $t0, $zero, 2
    bne      $a0, $t0, N3
    addi     $v0, $zero, 1
    j        fibend
```

Function Call Example 5 — Recursion (continued):

- The code for fibonacci using registers (continued):

```
N3:      # compute $t1 = fibonacci(N-1)
        addi    $a0, $a0, -1      # compute N - 1
        jal     fibonacci
        add     $t1, $v0, $zero   # save result1 in $t1

        # compute $t2 = fibonacci(N-2)
        lw      $a0, -12($fp)
        # save $t1 on the stack first
        # grow stack temporarily
        addiu   $sp, $sp, -4
        sw      $t1, 0($sp)
        addi    $a0, $a0, -2      # compute N - 2
        jal     fibonacci
        add     $t2, $v0, $zero   # save result2 in $t2
        # get $t1 off the stack and shrink the stack
        lw      $t1, 0($sp)
        addiu   $sp, $sp, 4

        add     $v0, $t1, $t2     # compute answer = result1 + result2

fibend:  # Function epilogue: restore stack & frame pointers and return
        lw      $a0, 8($sp)       # restore original value of $a0 for caller
        lw      $ra, 4($sp)       # get return address from stack
        lw      $fp, 0($sp)       # restore the caller's frame pointer
        addiu   $sp, $sp, 24      # restore the caller's stack pointer
```

the add right after main's jal fibonacci

Function Call Example 5 — Recursion (continued):

- The code for fibonacci using registers (continued):


```

N3:      # compute $t1 = fibonacci(N-1)
        addi    $a0, $a0, -1      # compute N - 1
        jal     fibonacci
        add     $t1, $v0, $zero   # save result1 in $t1
                                   fib(9) + fib(8)

        # compute $t2 = fibonacci(N-2)
        lw      $a0, -12($fp)
        # save $t1 on the stack first
        # grow stack temporarily
        addiu   $sp, $sp, -4
        sw      $t1, 0($sp)
        addi    $a0, $a0, -2      # compute N - 2
        jal     fibonacci
        add     $t2, $v0, $zero   # save result2 in $t2
        # get $t1 off the stack and shrink the stack
        lw      $t1, 0($sp)
        addiu   $sp, $sp, 4

        add     $v0, $t1, $t2     # compute answer = result1 + result2

fibend:  # Function epilogue: restore stack & frame pointers and return
        lw      $a0, 8($sp)      # restore original value of $a0 for caller
        lw      $ra, 4($sp)      # get return address from stack
        lw      $fp, 0($sp)      # restore the caller's frame pointer
        addiu   $sp, $sp, 24     # restore the caller's stack pointer
  
```

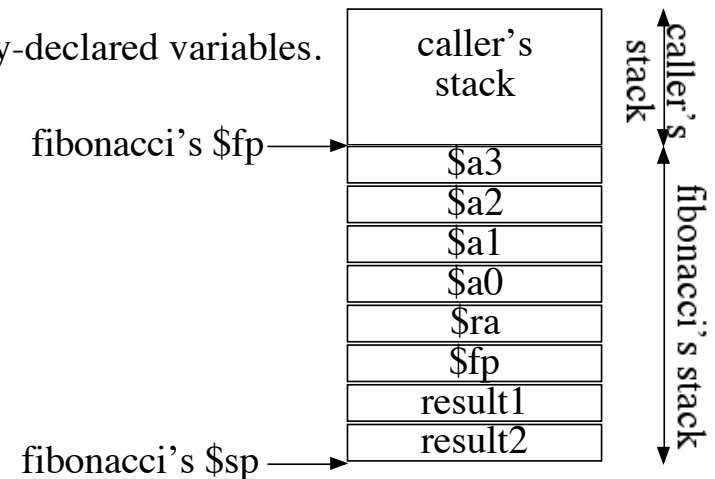


Function Call Example 5 — Recursion (continued):

- Using local variables:

- Basic idea: Create enough space on the stack initially to hold locally-declared variables.
- The C code would be:

```
int fibonacci( int N ) {  
    int result1;  
    int result2;  
    /* test for base cases not shown here... */  
    result1 = fibonacci( N - 1 );  
    result2 = fibonacci( N - 2 );  
    return result1 + result2;  
} /* fibonacci */
```



- For “local” variables (**result1** and **result2** in this case), create space on the stack.
 - Add enough space to the stack size, 8 bytes in this case.
 - Add extra space, if needed, to meet double-word aligned requirement (not needed this time).
 - Order of locals on the stack entirely up to the programmer — no convention for this.
 - No code in other functions will need to access this space.

Function Call Example 5 — Recursion (continued):

- The code for fibonacci using local variables:

```
fibonacci:
    # Prologue: set up stack and frame pointers for fibonacci
    # Need two local variables to hold the results of the two
    # recursive calls to fibonacci
    addiu    $sp, $sp, -32    # allocate stack space -- need 32 here
    sw       $fp, 8($sp)     # save caller's frame pointer
    sw       $ra, 12($sp)    # save return address
    sw       $a0, 16($sp)    # save parameter value
    addiu    $fp, $sp, 28    # setup fibonacci's frame pointer

    # if N == 1, return 1
    addi     $t0, $zero, 1
    bne      $a0, $t0, N2
    addi     $v0, $zero, 1
    j        fibend

N2:         # if N == 2, return 1
    addi     $t0, $zero, 2
    bne      $a0, $t0, N3
    addi     $v0, $zero, 1
    j        fibend
```

Function Call Example 5 — Recursion (continued):

- The code for fibonacci using local variables (continued):

```
N3:      # compute result1 = fibonacci(N-1)
        addi    $a0, $a0, -1      # compute N - 1
        jal     fibonacci
        sw      $v0, 4($sp)       # save result1

        # compute result2 = fibonacci(N-2)
        lw      $a0, 16($sp)
        addi    $a0, $a0, -2      # compute N - 2
        jal     fibonacci
        sw      $v0, 0($sp)       # save result2

        lw      $t1, 4($sp)       # $t1 = result1
        lw      $t2, 0($sp)       # $t2 = result2
        add     $v0, $t1, $t2     # compute answer = result1 + result2

fibend:  # Function epilogue: restore stack & frame pointers and return
        lw      $a0, 16($sp)      # restore original value of $a0 for caller
        lw      $ra, 12($sp)      # get return address from stack
        lw      $fp, 8($sp)       # restore the caller's frame pointer
        addiu   $sp, $sp, 32      # restore the caller's stack pointer
        jr      $ra              # return to caller's code
```


Function Call Example 5 — Recursion (continued):

- Summary of choice between registers and local variables:
 - Using t registers (the registers choice):
 - Create space on the stack just before the next call to fibonacci.
 - Remove the space from the stack just after the next call to fibonacci.
 - Using local variables:
 - Creates space on the stack for the variables at the beginning (during the prologue) of the function.
 - Removes that space at the end (during the epilogue) of the function.
- These two example programs are available as:

fibonacci-register.s

fibonacci-local.s

Memory Hierarchy

Read: Sections 5.1 to 5.3 (4th edition).

- Users want:
 - Lots of fast (quick response) memory.
 - Low cost.
- Solution: Hybrid systems.
 - Memory hierarchy containing different types of memory.


Memory Types:

- Fast memory:
 - SRAM — static random access memory.
 - Value stored on a pair of inverting gates; need 6 transistors per bit.
 - Value remains as long as power is supplied to the memory (hence *static*).
 - Fast: 0.5 to 2.5 ns (nanosecond) access time.
 - Expensive: \$2K - \$5K per GigaByte (2008, fm page 453).

Memory Types (continued):

- Not so fast memory — This is what computer manufacturers mean when they advertise memory in a computer.
- DRAM — dynamic RAM.
- Value is stored in a capacitor (charged or not charged), 1 transistor per bit.
- Must be refreshed, “read” value about every 50 ms (milliseconds).
- Dense, many more bits on same size chip (compared to SRAM).
- Slow: 50 to 70 ns, 5 to 10 times slower than SRAM.
- Cheap: \$20 - \$75 per GigaByte (2008, fm page 453); \$12 per GigaByte (purchase made in August 2013).
- DRAM variations:
 - SDRAM — synchronous dynamic RAM.
 - Uses data input register and data output register to buffer data.
 - 3 clock cycles to get first word.
 - 1 clock cycle per word for successive words.
 - Processor does not have to take into account delay, clock does that for it.
 - DDR — double data rate RAM.
 - Read (or write) a value on both the leading and trailing clock edges.

Memory Hierarchy:



Locality:


- A principle that makes having a memory hierarchy a good idea.
- If an item is referenced:
 - Temporal locality: The item will tend to be referenced again soon.
 - Spatial locality: Nearby items will tend to be referenced soon.
- Why does code have locality?
- Why does data have locality?
- Our initial focus:
 - Two levels of memory: upper and lower.
 - Block: minimum unit of memory.
 - Hit: data requested is in the upper level of memory.
 - Miss: data requested is not in the upper level of memory.

Cache — Upper Memory:

- Closest memory to the CPU.
- Two issues:
 - How do we know if a data item is in the cache?
 - If it is in the cache, how do we find it?
- Our first example:
 - Block size is one word of data; 4 bytes.
 - “Direct Mapped”
 - For each block of data at the lower level, there is exactly one location in the cache where it might be.
 - E.g., lots of items at the lower level “share” one location in the upper level.

Direct Mapped Cache:



- Mapping: an address is modulo the number of blocks in the cache.
- E.g., an 8 block cache for a 32 block memory:



- Cache location taken from the 3 least significant bits of the memory address, since $2^3 = 8$.
- Cache size is always a power of 2 for this reason!

Direct Mapped Cache (continued):

- Another example: a 16-block cache for a 64 block memory:



- Here, we need 4 bits for the cache address — take the 4 least significant bits of the address.
- For comparison: in the 8 block cache, the red memory locations mapped to the gray cache, and green to blue.


Direct Mapped Cache (continued):

- Issues:

- How do we know which value from the lower memory is currently in the cache location?
 - Store “tag” in the cache, using the upper part of the memory address (part that is not the cache address).
- How do we know if the value in the cache is a valid value?

```

0 0 0 0 0 0 0 1 1 1 1 1 1 1
0 0 0 0 1 1 1 1 0 0 0 0 1 1
0 0 1 1 0 0 1 1 0 0 1 1 0 0
0 1 0 1 0 1 0 1 0 1 0 1 0 1
    
```



```


0 0 0 0 0 0 0 0 0 1 0 0 0 0 0 0
0 0 0 0 0 0 1 0 0 0 0 0 0 0 0 0
0 1 0 0 0 1 0 0 0 0 1 0 0 0 1 0
    
```

Tag field

Valid bit

- Use “valid” bits.
- Set all valid bits to zero when program starts.
- Also, when a “context switch” occurs.

etc.



```

0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 1 1 1 1 0 0 0 0 1 1 1 1
0 0 1 1 0 0 1 1 0 0 1 1 0 0 1 1
0 1 0 1 0 1 0 1 0 1 0 1 0 1 0 1
    
```


CSs 252 — Computer Organization

17

8 — Memory Hierarchy


Direct Mapped Cache (continued):

- MIPS Example:
 - Block size is one word, 4 bytes.
 - Need 2 bits for the Byte offset.
 - Cache contains 1,024 blocks.
 - Need 10 bits for the Index.
 - $32 \text{ bits} - 10 \text{ for Index} - 2 \text{ for Byte} = 20 \text{ bits for the Tag}.$
- Issues:
 - How do we know which value from the lower memory is currently in the cache location?
 - Store the tag in the cache.
 - How do we know if the value in the cache is a valid value?
 - Valid bit in the cache.
- Cache width = $32 \text{ bits data} + 20 \text{ bits Tag} + 1 \text{ bit Valid} = 53 \text{ bits}.$
- What kind of locality is this?



Direct Mapped Cache (continued):

- Spatial locality.



Direct Mapped Cache (continued):

- Calculations:
 - Block size is 4 words, 16 bytes.
 - 2 bits for the Byte offset.
 - 2 bits for the Word offset.
 - Cache contains 4,096 blocks (rows).
 - 12 bits for the Index.
 - $32 \text{ bits} - 12 \text{ bits for Index} - 2 \text{ bits for Word} - 2 \text{ bits for Byte} = 16 \text{ bits for the Tag}.$
- Issues:
 - How do we know which value from the lower memory is currently in the cache location?
 - Store tag in the cache (same answer).
 - How we know if the value in the cache is a valid value?
 - Valid bit in the cache (same answer).
- Cache width = $4 * (32 \text{ bits of data per word}) + 16 \text{ bits Tag} + 1 \text{ bit Valid} = 145 \text{ bits}.$
- The block offset determines which word passes the multiplexor.

Hits vs. Misses:

- Read hits.
 - This is what we want!
- Read misses.
 - Stall the CPU.
 - Fetch block from memory.
 - Deliver block to the cache.
 - Restart the CPU.
- Write hits:
 - Can replace data in the cache and memory (*write-through*).
 - Write the data only into the cache (*write-back* the cache later).
- Write misses:
 - Read the entire block into the cache.
 - Then write the word into the cache.
 - Then, replace data in memory when writing to cache (*write-through*), or later (*write-back*).

Split Cache:

- Most systems use a *split* cache:
 - Usually for the Level 1 cache (the one closest to the CPU).
- Using one cache (instead of a split cache) allows the sharing of the cache resource:
 - The space in the cache can be applied to code or data, as needed for individual programs.
 - But:
 - Code tends to exhibit strong temporal locality.
 - And, also has spatial locality.
 - Data tends to exhibit strong spatial locality.
- Splitting the cache allows the data cache to have spatial locality.

Associativity:

- Can reduce the miss ratio of a cache by using associativity.
- Allows multiple locations in the cache where the contents of a particular memory location might reside.
- Can have 2-way, 3-way, 4-way, etc., associativity.
 - Note: 1-way set associative == direct mapped.
- Consider an array of integers where we want to process every other element.
 - Direct mapped cache can hold only 4 values from the array. Other 4 elements of the cache are unused.
 - 2-way set associative allows 8 elements of the array to be in the cache at once. All elements of the cache are utilized.
 - Decreases the miss ratio!

1-way set associative
(direct mapped)

Block	Tag	Data
0		
1		
2		
3		
4		
5		
6		
7		

Each memory location can map to only one spot in the cache.

2-way set associative

Set	Tag	Data	Tag	Data
0				
1				
2				
3				

Each memory location can map to two possible spots in the cache.

Associativity (continued):

- More possibilities:

4-way set associative

Set	Tag	Data	Tag	Data	Tag	Data	Tag	Data
0								
1								
2								
3								

Each memory location can map to four possible spots in the cache.

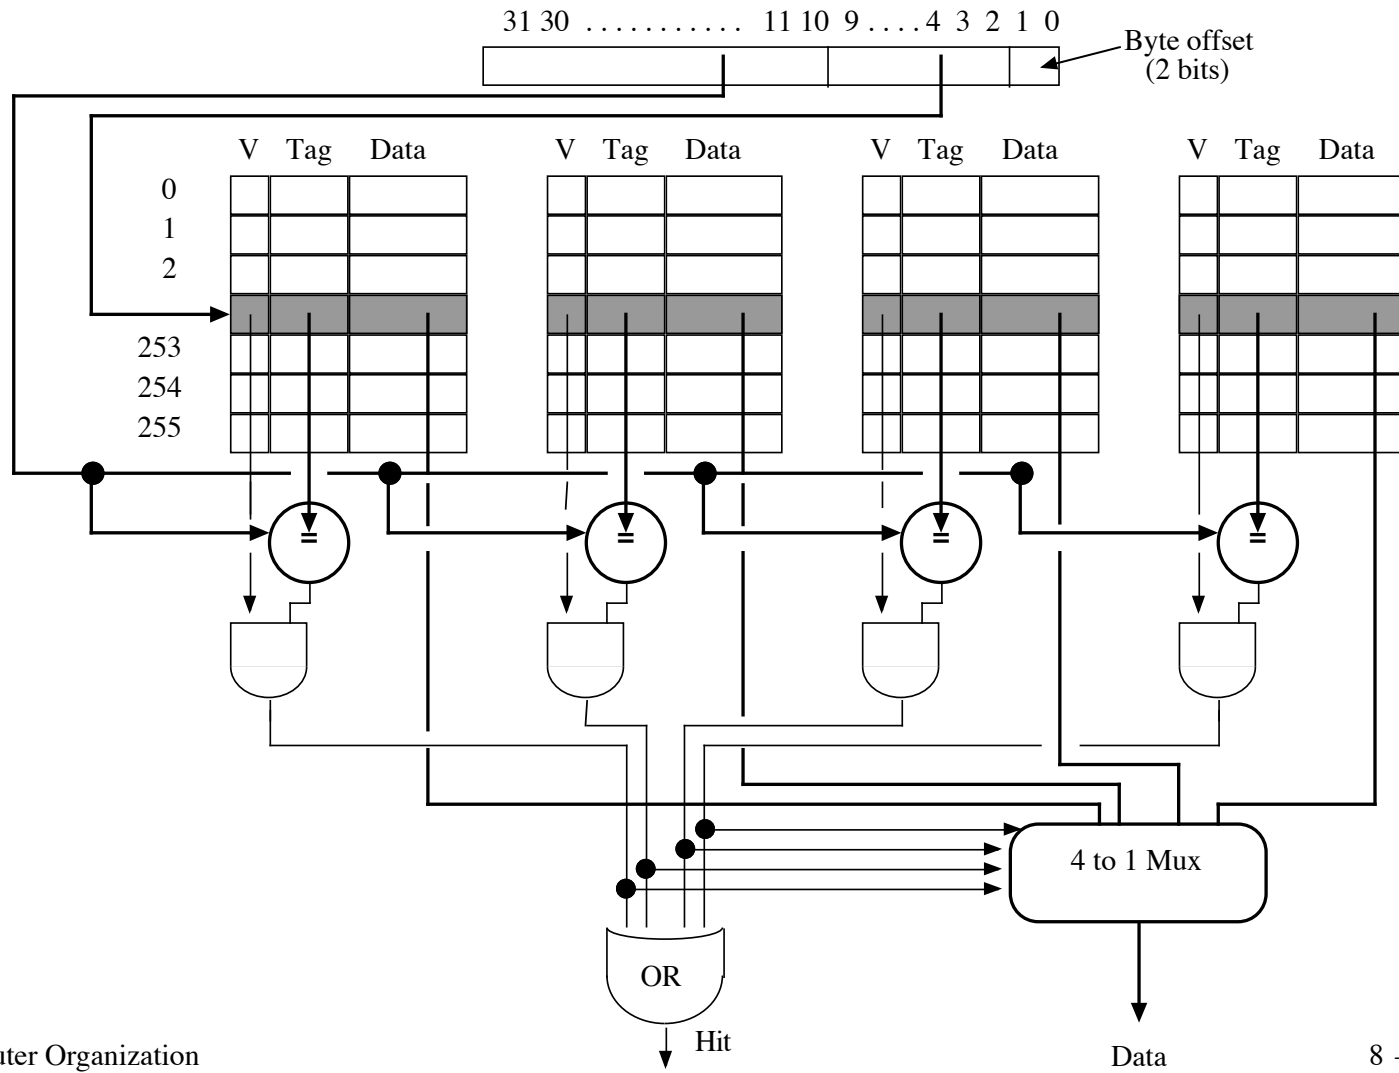
8-way set associative

	Tag	Data	Tag	Data	Tag	Data	Tag	Data	Tag	Data	Tag	Data	Tag	Data	Tag	Data
0																
1																
2																
3																

Each memory location can map to eight possible spots in the cache.

Associativity (continued):

- An implementation of a 4-way set associative cache:

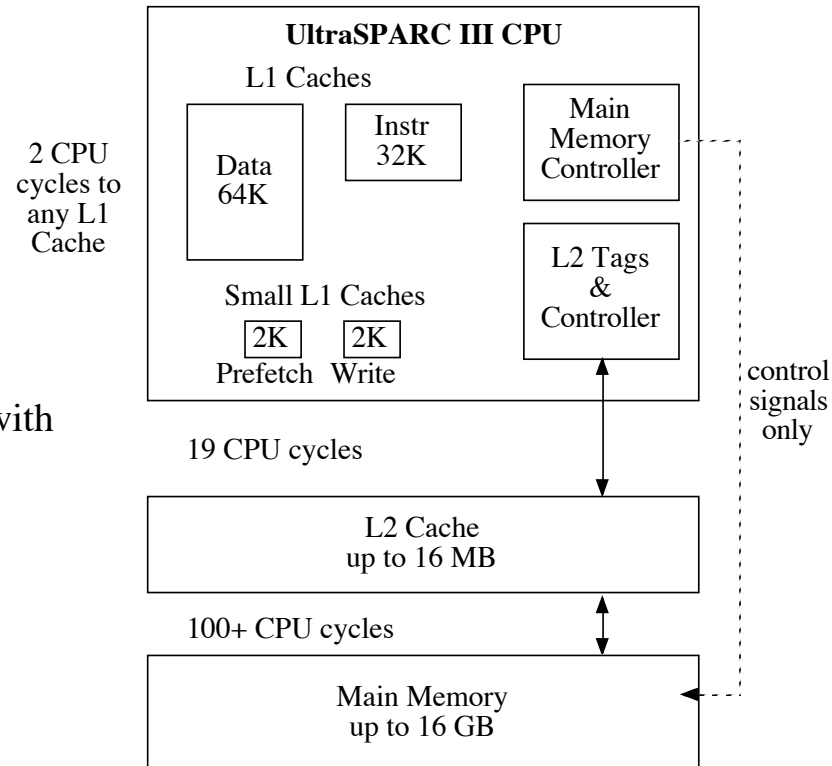


Decreasing miss penalty:

- Add a second level cache:
 - Often the primary cache is on the same chip as the processor.
 - Secondary (level-2) cache is off-chip.
- For dual-core (and multi-core) designs:
 - The primary (level-1) cache is with the core
 - 1 instruction memory cache per core.
 - 1 data memory cache per core.
 - The secondary (level-2) cache is on the chip and shared by all the cores.
 - There may (or may not) be a third (level-3) cache. This would be off the chip.

Sun Example:

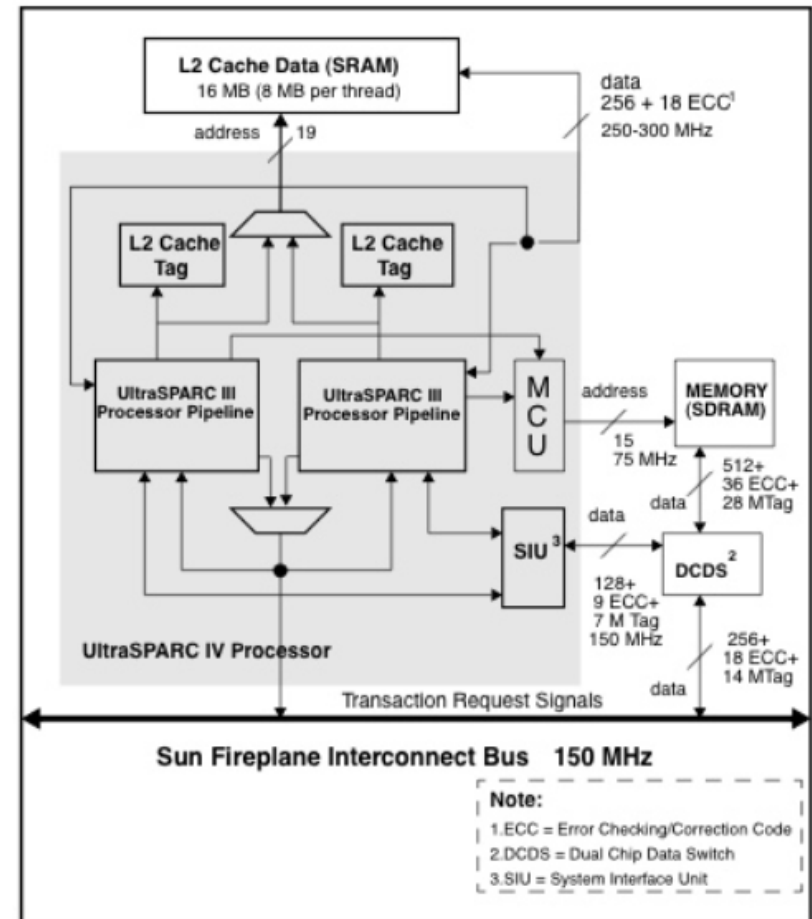
- Sun UltraSPARC III CPU.
 - 32 byte (256 bit) dedicated data path for L2 cache.
 - 128 bit data path to System (memory, I/O, any remote CPUs)
 - Runs at 1/8 of the CPU's clock speed.
 - 2.4 GB/sec transfer rate.
 - Memory controller: up to 15 outstanding load/store requests, with out-of-order completion.
 - Cache tags for L2 on chip to support cache coherency and snooping.
 - System interface on each chip (not shown in diagram)
 - Connects to System interconnect.
 - Connects to I/O and other CPUs.
 - 29 million transistors.
 - 1368 pins



Sun Example (continued):

- Sun UltraSPARC IV: available February 2004
 - Two UltraSPARC III processor cores on a single chip.
 - 1369 pins (almost pin-compatible).
 - Each core has its own L1 Data and L1 Instruction cache.
 - L2 Cache not on the chip.
 - L2 Tags are on the chip; each core has its own copy.

Figure 1-1 Basic UltraSPARC IV Processor



Sun Example (continued):

- Sun UltraSPARC IV+: available Oct 2005.
 - Each core has L1 data and L1 instruction cache.
 - L1 instruction cache: 64 KB, 64-byte line size.
 - L1 data cache: 64 KB (same as before). Uses a “write-through” policy to maintain cache coherency.
 - Chip has on-board L2 Cache, both Tag and Data.
 - Shared by the two cores.
 - Reduced from 16 MB to 2MB.
 - One read or write request every 2 clock cycles.
 - Uses “copy-back” policy on writes.
- Taken from: UltraSPARC IV+ Processor, User's Manual Supplement, Version 1.0, October 2005.

